

A new method to simulate the Bingham and related distributions in directional data analysis with applications

John T. Kent, Asaad M. Ganeiber and Kanti V. Mardia
Department of Statistics, University of Leeds

August 22, 2013

Abstract

A new acceptance-rejection method is proposed and investigated for the Bingham distribution on the sphere using the angular central Gaussian distribution as an envelope. It is shown to have high efficiency and to be straightforward to use. The method can also be extended to Fisher and Fisher-Bingham distributions on spheres and related manifolds.

Some key words: acceptance-rejection, angular central Gaussian distribution, Fisher distribution, Grassmann manifold, Kent distribution, special orthogonal group, Stiefel manifold, von Mises distribution.

1 Introduction

Directional data analysis is concerned with statistical analysis on various non-Euclidean manifolds, starting with circle and the sphere, and extending to related manifolds. Comprehensive monographs are available for statistical analysis in this setting; see, e.g., Chikuse (2003); Fisher et al. (1987); Mardia and Jupp (2000). However, the subject of simulation has received much less coverage, with the key contributions scattered through the literature.

The need for effective simulation methods has grown in recent years as directional distributions have become components in more sophisticated statistical models, which are studied using MCMC methods. For example, Green and Mardia (2006) used the matrix Fisher distribution for random 3×3 rotation matrices in a Bayesian model to align two unlabelled configurations of points in \mathbb{R}^3 , with an application to a problem of protein alignment in bioinformatics.

In general there are suitable direct methods of simulation, especially methods based acceptance rejection, for the simpler directional models. However, it is necessary to resort to cumbersome MCMC methods for the more complicated distributions. The purpose of this paper is to extend availability of acceptance rejection methods to a wider class of directional distributions. The starting point is a new acceptance rejection method for the Bingham distribution, which can then be used as a building block in a wider range of applications.

The paper is organized as follows. Following some background and preparation in Section 2, the new acceptance rejection simulation method for the Bingham distribution is proposed and analyzed in Section 3. Special cases and extensions are covered in Sections 4 and 5. Finally Section 6 sets the results of this paper in context by reviewing the literature and summarizing the best available methods in different settings.

The unit sphere $S_p = \{x \in \mathbb{R}^q : x^T x = 1\}$, $p \geq 1$, comprises the unit vectors in \mathbb{R}^q , where throughout the paper p and q are related by $q = p + 1$. The surface area of S_p is given by $\pi_q = 2\pi^{q/2}/\Gamma(q/2)$ and the differential element of surface area can be written as $[dx]$. Thus the uniform distribution on S_p can be written as $\pi_q^{-1}[dx]$. A more explicit formula can be given using polar coordinates. For example, the circle S_1 can be parameterized by $\theta \in [0, 2\pi)$ with uniform measure $d\theta/(2\pi)$. The sphere S_2 can be parameterized by colatitude $\theta \in [0, \pi]$ and longitude $[0, 2\pi)$ with uniform measure

$$\sin \theta d\theta d\phi/(4\pi). \quad (1.1)$$

Strictly speaking a probability density on a manifold is a density with respect to an underlying measure. In Euclidean space \mathbb{R}^p the underlying measure is usually taken to be Lebesgue measure dx without explicit comment. But on other manifolds more care is needed. This paper is concerned with spheres and related compact manifolds for which there is a natural underlying uniform measure with a finite total measure. To avoid repeated occurrences of normalizing constants such as π_q and differential elements such as $[dx]$, all such probability densities will be expressed with respect to the uniform distribution. Thus we will write the density for the uniform distribution on S_2 as $f(x) = 1$ (with respect to itself) rather than as $f(x) = 1/(4\pi)$ (with respect to $[dx]$) or as $f(x) = \sin \theta$ (with respect to $d\theta d\phi$).

2 Background

Recall the acceptance-rejection method of simulation. Consider two densities,

$$f(x) = c_f f^*(x), \quad g(x) = c_g g^*(x) \quad (2.1)$$

where f^* and g^* are known functions, but where the normalizing constants may or may not have a known explicit form. Suppose it is possible to simulate easily from g and it is desired to simulate observations from f . The key requirement is that there is a known bound of the form

$$f^*(x) \leq M^* g^*(x) \text{ for all } x \quad (2.2)$$

for some constant M^* . The acceptance-rejection algorithm proceeds as follows.

Step 1. Simulate $X \sim g$ independently of $W \sim \text{Unif}(0, 1)$.

Step 2. If $W < f^*(X)/\{M^*g^*(X)\}$, then accept X .

Step 3. Otherwise go back to step 1.

Comments

- (a) If we set $M = c_f M^* / c_g$, then (2.2) can be expressed equivalently as $f(x) \leq M g(x)$ for all x .
- (b) The bound M satisfies $M \geq 1$. The number of trials needed from g is geometrically distributed with mean $M \geq 1$. The efficiency is defined by $1/M$. For high efficiency the bound M should be as close to 1 as possible.
- (c) The algorithm can be used even if the normalizing constants do not have a known explicit form. However, to compute the efficiency analytically, it is necessary to know the normalizing constants.
- (d) Suppose the density $g(x) = g(x; b)$ depends on a parameter b with corresponding bound $M^*(b)$ in (2.2). If the normalizing constant $c_g = c_g(b)$ has a known explicit form, then it is possible to maximize the efficiency with respect to b , even if c_f does not have a known explicit form.

When developing acceptance-rejection simulation methods for directional distributions, there are several issues to consider:

- the need for good efficiency for a wide range of concentration parameters for f , ranging from uniform to highly concentrated. In similar problems on \mathbb{R}^p , the task is simpler when distributions are closed under affine transformations; in such cases it is sufficient to consider just a single standardized form of the distribution for f .
- the challenge in finding a tractable envelope distribution.
- the presence of trigonometric factors in the base measure when expressed in polar coordinates, such as in (1.1).

Next, using the concavity of the log function, we give a general inequality which is useful in the construction of acceptance-rejection algorithms. Consider the function of $u \geq 0$,

$$\phi(u) = \frac{q}{2} \log(1 + 2u/b) - u - \frac{q}{2} \log(1 + 2u_0/b) + u_0, \quad (2.3)$$

where $q > 0$ and $0 < b < q$ are fixed constants and $u_0 = (q - b)/2$. The last two terms on the righthand side of (2.3) are constants, chosen so that $\phi(u_0) = 0$. The value of u_0 is chosen so that the function $\frac{q}{2} \log(1 + 2u/b)$ has slope 1 at $u = u_0$; hence $\phi'(u_0) = 0$. Also note that $\phi''(u) < 0$ for $u \geq 0$ so that $\phi(u)$ is a concave function. Therefore, $\phi(u) \leq 0$ for all $u \geq 0$. After exponentiating, this inequality can be re-arranged as

$$e^{-u} \leq e^{-(q-b)/2} \left(\frac{q/b}{1 + 2u/b} \right)^{q/2}. \quad (2.4)$$

p	1	2	3	4	5	10	50	100
eff.	66%	52%	45%	40%	36%	26%	12%	9%

Table 1: Efficiency of A/R simulation method for the multivariate normal distribution in p dimensions, using a multivariate Cauchy envelope.

To illustrate the usefulness of (2.4), we start in Euclidean space \mathbb{R}^p , $p \geq 1$, and construct an acceptance-rejection algorithm for the multivariate normal distribution using a multivariate Cauchy envelope. The multivariate normal distribution $N_p(0, \Sigma)$ has density

$$f(x) = c_f f^*(x), \quad f^*(x) = \exp\left(-\frac{1}{2}x^T \Sigma^{-1}x\right), \quad c_f = c_f(\Sigma) = |2\pi\Sigma|^{-1/2}, \quad (2.5)$$

for $x \in \mathbb{R}^p$. The multivariate Cauchy distribution $C_p(0, \Psi)$ has density

$$g(x) = c_g g^*(x), \quad g^*(x) = (1 + x^T \Psi^{-1}x)^{-q/2}, \quad c_g = c_g(\Psi) = \frac{\Gamma(q/2)}{\pi^{q/2}} |\Psi|^{-1/2}, \quad (2.6)$$

where for convenience we have substituted $q = p + 1$.

If we set $\Psi = b\Sigma$ so that the scatter matrix for the Cauchy is a scalar multiple of the covariance matrix for the normal, and if we set $u = \frac{1}{2}x^T \Sigma^{-1}x$, then the inequality (2.4) leads to a bound on the densities with

$$M(b) = 2^{-p/2} q^{q/2} e^{-q/2} \pi^{1/2} b^{-1/2} e^{b/2} / \Gamma(q/2).$$

Minimizing over $0 < b < q$ yields the optimal parameter $b = 1$ with optimal bound

$$M = M(1) = \sqrt{2\pi e} \left(\frac{q}{2e}\right)^{q/2} / \Gamma(q/2), \quad q = p + 1. \quad (2.7)$$

Table 1 gives collection of the efficiencies $1/M$ as a function of dimension p . For large p , $M \sim \sqrt{qe/2}$ by Stirling's formula.

Note that the efficiency declines slowly with the dimension, but is still high enough to be feasible even for dimension $p = 100$. Of course, this is just a toy example since there are better ways to simulate the normal distribution. However, it is important for the next section, both as a motivating example and as a limiting case.

3 Simulating the Bingham distribution with an angular central gaussian envelope

In this section we describe the ‘‘BACG’’ acceptance rejection method to simulate the Bingham distribution using the angular central Gaussian distribution as an envelope. As before, let $q = p + 1$.

The Bingham distribution, $\text{Bing}_p(A)$ on S_p , $p \geq 1$, where the parameter matrix A is $q \times q$ symmetric, has density

$$f_{\text{Bing}}(x) = c_{\text{Bing}} f_{\text{Bing}}^*(x), \quad f_{\text{Bing}}^*(x) = \exp(-x^T A x). \quad (3.1)$$

The normalizing constant $c_{\text{Bing}} = c_{\text{Bing}}(A)$ can be expressed as a hypergeometric function of matrix argument (Mardia and Jupp, 2000, p. 182), but is not sufficiently tractable to be of interest here. The use of a minus sign in the exponent is unconventional but simplifies later formulae. Since A and $A + cI$ define the same distribution for any real constant c , we may assume without loss of generality that the eigenvalues of A satisfy

$$0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_q. \quad (3.2)$$

The angular central Gaussian distribution, $\text{ACG}(\Omega)$ on S_p , where the parameter matrix Ω is $q \times q$ symmetric positive definite, takes the form

$$f_{\text{ACG}}(x) = c_{\text{ACG}} f_{\text{ACG}}^*(x), \quad f_{\text{ACG}}^*(x) = (x^T \Omega x)^{-q/2}, \quad c_{\text{ACG}} = |\Omega|^{1/2}. \quad (3.3)$$

The angular central Gaussian distribution is simple to simulate. If $y \sim N_q(0, \Sigma)$, where Σ is positive definite, then $x = y/\|y\| \sim \text{ACG}(\Omega)$ with $\Omega = \Sigma^{-1}$ (e.g., Mardia and Jupp, 2000, p. 182).

Setting $u = x^T A x$ in (2.4) and setting $\Omega = \Omega(b) = I + 2A/b$, $b > 0$, yields the envelope inequality on the starred densities

$$\begin{aligned} f_{\text{Bing}}^*(x) &= e^{-u} \\ &\leq e^{-(q-b)/2} \left(\frac{q/b}{1 + 2x^T A x/b} \right)^{q/2} \\ &= e^{-(q-b)/2} \left(\frac{q/b}{x^T \Omega x} \right)^{q/2} \\ &= e^{-(q-b)/2} (q/b)^{q/2} f_{\text{ACG}}^*(x), \end{aligned} \quad (3.4)$$

using the constraint $x^T x = 1$. The corresponding bound $M(b)$ takes the form

$$M(b) = c_{\text{Bing}} e^{-(q-b)/2} (q/b)^{q/2} |\Omega(b)|^{-1/2}. \quad (3.5)$$

It can be checked that $\log M(b)$ is convex in b with unique minimizing value given by the solution of

$$\sum_{i=1}^q \frac{1}{b + 2\lambda_i} = 1, \quad (3.6)$$

where the lefthand side of (3.6) ranges between ∞ and 0 as b ranges between 0 and ∞ . Let b_0 denote the solution to (3.6) and let $M(b_0)$ denote the optimal bound.

It does not seem possible to evaluate $M(b_0)$ in a more useful form analytically, but it is possible to say what happens asymptotically. Replace A in (3.1) by βA and think of A as a fixed matrix as $\beta > 0$ gets large. Provided the p largest eigenvalues of

λ_2	λ_3	Efficiency
0	0	100%
0	10	84%
10	10	58%
0	100	80%
100	100	53%

Table 2: Efficiency of the BACG A/R simulation method on S_p with $A = \text{diag}(0, \lambda_2, \lambda_3)$ for the Bingham distribution with an ACG envelope.

A are strictly positive, the ACG distribution (restricted to a hemisphere about the mode) converges to a p -dimensional multivariate Cauchy distribution, the Bingham distribution converges to a p -dimensional multivariate normal distribution, b_0 converges to 1 and the bound $M(b_0)$ converges to the bound (2.7) where the dimensions p and $q = p + 1$ have the same meanings in both sections.

Empirically, it has been noticed that the limiting case is the worst possible case. For smaller values of the concentration matrix A , the efficiencies will be higher. Table 2 illustrates the pattern for $p = 2$, i.e. $q = 3$. The efficiency is never lower than 52%, the value from Table 1 for $p = 2$. This limiting value is attained in the concentrated bipolar case (when $\lambda_2 = \lambda_3$ is large). The girdle case ($\lambda_2 = 0$) has higher efficiencies. Each entry in this table has been constructed from one million simulations, so that the standard errors are negligible.

4 Manifolds and models in directional data analysis

In order to prepare for special cases and extensions of the Bingham distribution, it is helpful to give a brief survey of some of the various manifolds and models used in directional data analysis. For each of these manifolds there is a unique invariant measure which can be used to define a uniform distribution.

4.1 The sphere S_p revisited

A general model on the sphere S_p is the Fisher-Bingham distribution with density

$$f_{\text{FB}}^*(x) = \exp(\kappa x^T \mu_0 - x^T A x) \quad (4.1)$$

where $\kappa \geq 0$, $\mu_0 \in S_p$ and $A(q \times q)$ is symmetric, without loss of generality with smallest eigenvalue equal to 0. If $A = 0$ the model reduces to the von Mises ($p = 1$), the Fisher ($p = 2$), or the von Mises-Fisher (any $p \geq 1$) distribution. If $\kappa = 0$, the model reduces to the Bingham distribution considered in Section 3.

4.2 Real projective space $\mathbb{R}P_p$

Real projective space is defined as the quotient space $\mathbb{R}P_p = S_p / \{1, -1\}$ in which two antipodal points or “directions” $\pm x$ are identified with one another to represent the same “axis”. Since the Bingham and ACG densities have the property of antipodal symmetry, $f(x) = f(-x)$, they can also be viewed as densities on $\mathbb{R}P_p$.

4.3 Complex projective space $\mathbb{C}P_p$

Another quotient space of the sphere is complex projective space, $\mathbb{C}P_p = S_{2p+1} / S_1$. To understand this space, suppose a unit vector $x \in \mathbb{R}^{2q}$, $q = p + 1$, is partitioned as $x^T = (x_1^T, x_2^T)$ where x_1 and x_2 are q -dimensional. The information in x can also be represented by a q -dimensional complex vector $z = x_1 + ix_2$. Then $\mathbb{C}P_p$ is obtained from S_{2p+1} by identifying the scalar multiples $e^{i\theta}z$ with one another for all $\theta \in [0, 2\pi)$.

If the $2q \times 2q$ symmetric concentration matrix A for a Bing_{2p+1} distribution can be partitioned in the form

$$A = \begin{bmatrix} A_1 & -A_2 \\ A_2 & A_1 \end{bmatrix},$$

where A_1 is symmetric and A_2 is skew symmetric, then the quadratic form $-x^T A x$ in the exponent of the Bingham density can be expressed in complex notation as $-z^* A_C z$ where $A_C = A_1 + iA_2$. In terms of z , the density possesses complex symmetry, $f(z) = f(e^{i\theta}z)$ for all $\theta \in [0, 2\pi)$. When expressed in complex notation this distribution is known as the complex Bingham distribution $\text{CB}_p(A_C)$; it can also be viewed as a distribution on $\mathbb{C}P_p$ (Kent, 1994).

4.4 The special orthogonal group $\text{SO}(r)$

The special orthogonal group $\text{SO}(r)$ is the space of $r \times r$ rotation matrices, $\text{SO}(r) = \{X \in \mathbb{R}^{r \times r} : X^T X = I_r, |X| = 1\}$. A natural parametric distribution is given by the matrix Fisher distribution $MF_r(F)$, with $r \times r$ parameter matrix F . The density is given by

$$f^*(X) = \exp\{\text{tr}(F^T X)\}. \quad (4.2)$$

To describe the concentration properties of this distribution, it is helpful to give F a signed singular value decomposition

$$F = U \Delta V^T. \quad (4.3)$$

The adjective “signed” means that U and V are $r \times r$ rotation matrices and the elements of the diagonal matrix Δ satisfy $\delta_1 \geq \dots \geq \delta_{r-1} \geq |\delta_r|$, where the final element is negative if and only if $|F| < 0$.

4.5 The Stiefel manifold $\mathcal{V}_{r,q}$

Let $1 \leq r \leq q$ and define the Stiefel manifold $V_{r,q} = \{X_1 \in \mathbb{R}^{q \times r} : X_1^T X_1 = I_r\}$ to be the space of $q \times r$ column orthonormal matrices X_1 , say.

The matrix Fisher-Bingham distribution on $V_{r,q}$, denoted $MFB(F_1, A, C)$, with parameter matrices $F_1(q \times r)$, $A(q \times q)$ symmetric and $C(q \times q)$ symmetric, is defined by the density

$$f_{\text{MFB}}(X) \propto \text{etr}(F_1^T X_1 - C X_1^T A X_1). \quad (4.4)$$

Special cases include the matrix Fisher distribution, denoted $MF(F_1)$, with density

$$f_{\text{MF}}(X) \propto \text{etr}(F_1^T X_1), \quad (4.5)$$

the “full” matrix Bingham distribution, denoted $MB(A, C)$, with density

$$f_{\text{MB}}(X) \propto \text{etr}(-C X_1^T A X_1), \quad (4.6)$$

and the “balanced” matrix Bingham distribution, denoted $MB(A)$, with density

$$f_{\text{MB-bal}}(X) \propto \text{etr}(-X_1^T A X_1). \quad (4.7)$$

If $r = q$ the balanced matrix Bingham distribution reduces to the uniform distribution. The reason is that X_1 is an orthogonal matrix in this case, so that $X_1 X_1^T = I_q$ and $\text{tr}(X_1^T A X_1) = \text{tr}(X_1 X_1^T A) = \text{tr}(A)$ is constant in (4.7).

If $r = q - 1$, it is possible to extend X_1 uniquely by adding an extra column to form a $q \times q$ rotation matrix X , say. Hence $V_{q-1,q}$ can be identified with $SO(q)$. However, the version of the matrix Fisher distribution in (4.2), with an $r \times r$ parameter matrix F , is more general than that in (4.5), with an $r \times (r - 1)$ parameter matrix F_1 .

If $r = q$, then $V_{q,q}$ is the same as the orthogonal group $O(q)$, which is twice the size of the special orthogonal group $SO(q)$, since in $O(q)$ a matrix X is allowed to have determinant ± 1 . Although the two densities (4.2) and (4.5) formally look the same, they live on different spaces.

4.6 The Grassmann manifold $\mathcal{G}_{r,q}$

Let $1 \leq r < q$. The Grassmann manifold $\mathcal{G}_{r,q}$ is defined to be the set of all r -dimensional subspaces of \mathbb{R}^q . It can be described as a quotient space of a Stiefel manifold $\mathcal{G}_{r,q} = V_{r,q}/O(r)$, in which a $q \times r$ column orthonormal matrix X_1 is identified with $X_1 R$ for all $r \times r$ orthogonal matrices R . It should be noted that the notation for this manifold is not standardized; some authors write $\mathcal{G}_{r,q-r}$ instead of $\mathcal{G}_{r,q}$.

Since $\text{tr}(X_1^T A X_1) = \text{tr}(R^T X_1^T A X_1 R)$, the balanced matrix Bingham distribution (4.7) on the Stiefel manifold $\mathcal{V}_{r,q}$ can also be viewed as a distribution on the Grassmannian manifold $\mathcal{G}_{r,q}$.

For every r -dimensional subspace in \mathbb{R}^q , there is a unique complementary $(q - r)$ -dimensional subspace. If X_1 and X_2 are column orthonormal matrices, whose columns are bases of these subspaces, then $X = (X_1 \ X_2)$ is a $q \times q$ orthogonal matrix.

Further X_1 follows a balanced matrix Bingham distribution on $\mathcal{G}_{r,q}$ with parameter matrix A if and only if X_2 follows a balanced matrix Bingham distribution on $\mathcal{G}_{q-r,q}$ with parameter matrix $-A$ (but be warned that the eigenvalues of $-A$ will not have the standardized form in (3.2)). Hence for simulation purposes, we may without loss of generality suppose that $r \leq q/2$.

5 Accidental isomorphisms

The two quotient manifolds $\mathbb{R}P_p$ and $\mathbb{C}P_p$ of direct interest for the Bingham distribution due to the existence of “accidental isomorphisms” in which the quotient manifold becomes identified with another familiar manifold through a quadratic mapping. These isomorphisms are called “accidental” because there does not seem to be any systematic pattern. In each case the uniform distribution on the quotient manifold maps to the uniform distribution on the new manifold, and the Bingham distribution maps to a distribution related to the von Mises-Fisher distribution on the new manifold. The implications for simulation are laid out in the next subsections.

5.1 $\mathbb{R}P_1 = S_1$

Euclidean coordinates on the circle can be represented in polar coordinates by $x = (x_1, x_2)^T$ where $x_1 = \cos \theta, x_2 = \sin \theta, \theta \in [0, 2\pi)$. Consider a two-to-one map to a new circle defined by $\phi = 2\theta$, with Euclidean coordinates $y = (y_1, y_2)^T$ where $y_1 = \cos \phi = x_1^2 - x_2^2, y_2 = \sin \phi = 2x_1x_2$. Note that the antipodal directions $\theta, \theta + \pi$ map to the same value of ϕ , so that the map is in fact a one-to-one map between $\mathbb{R}P_1$ and S_1 . A quadratic form in x can be rewritten as

$$x^T A x = \frac{1}{2}(a_{11} - a_{22})y_1 + a_{12}y_2 + \frac{1}{2}(a_{11} + a_{22}),$$

which is a linear function of y . Hence a Bingham distribution, whose density is quadratic in x on $\mathbb{R}P_1$ can be identified with a von Mises distribution, whose density is linear in y , on S_1 .

Similarly, in the ACG density the quadratic form $x^T \Omega x$ becomes a linear function of y , so the density in y reduces to the wrapped Cauchy density (Mardia and Jupp, 2000, p. 52).

Suppose A is diagonal, $A = \text{diag}(0, \lambda)$. In this case the dominant axis of the Bingham distribution is the x_1 -axis. The corresponding von Mises density takes the form

$$f_{\text{VM}}(y) \propto \exp(\kappa y_1), \quad \kappa = \lambda/2,$$

so that the corresponding von Mises density has its mode in the y_1 -direction. The corresponding wrapped Cauchy density, with $\Omega = I + 2A/b$, takes the form,

$$f_{\text{WC}}(y) = \frac{(1 - \rho^2)}{1 + \rho^2 - 2\rho y_1}$$

where $\rho = (\beta - 1)/(\beta + 1)$ (Mardia and Jupp, 2000, p. 52).

Hence the simulation method for the Bingham distribution with an ACG envelope can be recast as a simulation for the von Mises distribution with a wrapped Cauchy envelope. It turns out that this latter method is identical to the proposal of Best and Fisher (1979), even up to the choice of the optimal tuning constant b .

5.2 $\mathbb{C}P_1 = S_2$

The complex projective space $\mathbb{C}P_{k-2}$ arises in the study of shape for configurations of k landmarks in the plane, and the identification with S_1 when $k = 3$ was used to visualize the shape space for triangles of landmarks (Kendall, 1984). Kent (1994) showed that the complex Bingham distribution on $\mathbb{C}P_1$ can be identified with the Fisher distribution on S_2 .

Motivated by this accidental isomorphism, Kent et al. (2006) developed a complex Bingham quartic (CBQ) distribution on $\mathbb{C}P_p$, $p \geq 1$. When $p = 1$, this distribution reduces to the FB_5 distribution. Ganeiber (2012) developed an effective and reasonably efficient simulation method for the CBQ distribution for $p > 1$. However, since the technique is not based on an angular central Gaussian envelope, details will not be given here.

5.3 $\mathbb{R}P_3 = SO(3)$

There is a quadratic mapping taking an unsigned 4-dimensional unit vector $\pm x$ to a 3×3 rotation matrix $X = M(x) = M(-x)$, say. More specifically

$$M(x) = \begin{bmatrix} x_1^2 + x_2^2 - x_3^2 - x_4^2 & -2(x_1x_4 - x_2x_3) & 2(x_1x_3 + x_2x_4) \\ 2(x_1x_4 + x_2x_3) & x_1^2 + x_3^2 - x_2^2 - x_4^2 & -2(x_1x_2 - x_3x_4) \\ -2(x_1x_3 - x_2x_4) & 2(x_1x_2 + x_3x_4) & x_1^2 + x_4^2 - x_2^2 - x_3^2 \end{bmatrix} \quad (5.1)$$

(Mardia and Jupp, 2000, p. 285). Further a random axis $\pm x$ on $\mathbb{R}P_3$ follows a Bingham distribution if and only if the corresponding random matrix $M(x)$ follows a matrix Fisher distribution. In particular, if $A = \Lambda$ is diagonal, then $F = \Delta$ in (4.3) is also diagonal with the parameters related by

$$\lambda_1 = 0, \quad \lambda_2 = 2(\delta_2 + \delta_3), \quad \lambda_3 = 2(\delta_1 + \delta_3), \quad \lambda_4 = 2(\delta_1 + \delta_2). \quad (5.2)$$

Kent et al. (2012) gives some further details.

A simple way to simulate a rotation matrix from the matrix Fisher distribution $MF_3(F)$ for a general parameter matrix F with signed singular value decomposition (4.3) is given as follows. Using the BACG method simulate x from $\text{Bing}_3(\Lambda)$ with Λ given by (5.2), and let $M(x)$ denote the corresponding rotation matrix using (5.1). Then $UM(x)V^T$ follows the matrix Fisher distribution $MF_3(F)$. From Table 1, the efficiency will be at least 45%.

6 Simulating the Fisher-Bingham distribution with an angular central gaussian envelope

The von Mises-Fisher density on S_p takes the form (4.1) with $A = 0$. The elementary inequality $(1 - y)^2 \geq 0$, with $y = x^T \mu_0$ can be re-arranged to give

$$\begin{aligned} f_F^*(x) &\leq \exp \left[(\kappa/2) \left\{ (x^T \mu_0)^2 + 1 \right\} \right] \\ &= \exp \left\{ \kappa - (\kappa/2) x^T A x \right\} \\ &= e^\kappa f_{\text{Bing}}^*(x), \end{aligned} \tag{6.1}$$

where $A = I_q - \mu_0 \mu_0^T$. Hence an acceptance rejection simulation method for the von Mises-Fisher distribution can be constructed using a Bingham envelope.

The two sides of (6.1) match when $x = \mu_0$ so that it is not possible to get a tighter bound. In relative terms, the two starred densities are maximally different when $x = -\mu_0$. This difference matters most when κ is large, when the efficiency of acceptance-rejection with a Bingham envelope drops to 50%; the efficiency rises to 1 as $\kappa \rightarrow 0$. Empirically the efficiency lies between these two extremes for intermediate values of κ .

The inequality (6.1) can be combined with Section 3 to provide a method to simulate the von Mises-Fisher distribution with an ACG envelope. Of course there is no need for a new method for the von Mises-Fisher distribution. Good methods are already available; see the Section 8 for a discussion. However, the bounds of this section can be combined with the previous section to simulate the Fisher-Bingham distribution with an ACG envelope.

The Fisher-Bingham distribution on S_p takes the form (4.1) and can be bounded by a Bingham density

$$f_{\text{FB}}^*(x) \leq \exp(\kappa - x^T A^{(1)} x),$$

where $A^{(1)} = A + (\kappa/2)(I - \mu_0 \mu_0^T)$. Then Section 3 can be used to bound this Bingham density by an ACG density.

Say the Fisher-Bingham distribution is “aligned” if μ_0 is an eigenvector of A and if the density has its mode at $x = \mu_0$. In this case the Bingham envelope usually has an efficiency of at least 50%, with the efficiency falling below this level only when the density excessively flat at its mode and the concentration is high. Under high concentration this situation corresponds to the case where the limiting normal distribution would have a singular covariance matrix.

7 The balanced matrix Bingham distribution

The density for the balanced matrix Bingham distribution was given by (4.7). It can be viewed as either a density on the Stiefel manifold $\mathcal{V}_{r,q}$ (which is invariant under multiplication on the right by $r \times r$ orthogonal matrices, or on the Grassmannian manifold $\mathcal{G}_{r,q}$. The $q \times q$ concentration matrix A has the same form as for the Bingham distribution in Section 3.

The matrix ACG distribution, denoted $MACG_{r,q}(\Omega)$, where Ω is a positive definite symmetric $q \times q$ matrix, is also lies on $\mathcal{V}_{r,q}$. It is also invariant under rotation on the right and hence can also be viewed as a distribution on the Grassmann manifold $\mathcal{G}_{r,q}$. The density takes the form

$$g_{MACG}^*(X) = |X^T \Omega X|^{-q/2}, \quad c_g = |\Omega|^{r/2}$$

(e.g. Chikuse, 2003, p. 40). Simulations from this distribution can be constructed as follows. Let Y be a $q \times r$ matrix whose columns are independently normally distributed, $N_q(0, \Omega^{-1})$. Set $X = Y(Y^T Y)^{-1/2}$ using the symmetric square root of a positive definite matrix. Then $X \sim MACG_{r,q}(\Omega)$.

If Ω is related to A by $\Omega = \Omega(b) = I_q + 2A/b$ as in Section 3, then the balanced matrix Bingham density can be bounded by the matrix ACG density by using the inequality in (2.4) r times. Namely, let the eigenvalues of $X^T A X$ be denoted $u_1 \geq \dots \geq u_r \geq 0$. Since $f_{MB-bal}^*(X) = \exp(\sum u_i)$ and $f_{MACG}^*(X) = \{\prod(1 + 2u_i/b)\}^{-q/2}$, applying (2.4) r times yields the envelope bound

$$M(b) = c_{MB-bal} \{e^{-(q-b)/2} (q/b)^{q/2} |\Omega(b)|^{-1/2}\}^r. \quad (7.1)$$

Optimizing over b yields the same equation (3.6) as before with the same value for the optimal value b_0 .

The efficiency is expected to decline as r increases. However, as noted before, we may restrict attention to the case $r \leq q/2$. More numerical investigation is needed of the efficiency in this setting.

8 Review and commentary on different simulation methods

Since the simulation literature for directional distributions is widely scattered, it is useful to summarize the best simulation methods for various distributions of interest. Table 8 lists several common distributions on different spaces, together with the recommended method of simulation.

Recently, some MCMC simulation methods on manifolds have been proposed by Kume and Walker (2009) (Fisher-Bingham on S_p), Habeck (2009) (matrix Fisher on $SO(3)$), Hoff (2009) (matrix Fisher-Bingham distributions on Stiefel and Grassmann manifolds) and Byrne and Girolami (2013) (more general distributions). However, there is still an ongoing investigation into the efficiency of different MCMC methods, so that the table entry will just state ‘‘MCMC’’ when there is not a simpler more specific recommendation. Further details are given in the following subsections.

8.1 Uniform distribution on S_p

The simplest general method to simulate a uniform distribution on the unit sphere S_p is to set $x = u/||u||$ where $u \sim N_q(0, I_q)$, $q = p + 1$. In low dimensions

Table 3: Recommended simulation methods various distributions on different directional spaces

Distribution	Space	Simulation method
von Mises-Fisher	S_p	Wood (1987)
Bingham	S_p or $\mathbb{R}P_p$	BACG
complex Bingham	S_{2p+1} or $\mathbb{C}P_p$	Kent et al. (2004)
complex Bingham quartic	S_{2p+1} or $\mathbb{C}P_p$	Kent et al. (2006)
aligned Fisher-Bingham	S_p	BACG-based
general Fisher-Bingham	S_p	MCMC
matrix Fisher	$\mathcal{V}_{r,q}$	MCMC
matrix Fisher	$SO(3)$	BACG
matrix Fisher	$SO(p)$, $p > 3$	MCMC
balanced matrix Bingham	$\mathcal{V}_{r,q}$ or $\mathcal{G}_{r,q}$	BACG
general matrix Bingham	$\mathcal{V}_{r,q}$	MCMC
matrix Fisher-Bingham	$\mathcal{V}_{r,q}$	MCMC

there are sometimes simpler methods using polar coordinates. E.g. on the circle S_1 , let $\theta \sim \text{Unif}(0, 2\pi)$. On the sphere S_2 with colatitude θ and longitude ϕ , let $\cos \theta \sim \text{Unif}(-1, 1)$ independently of $\phi \sim \text{Unif}(0, 2\pi)$.

On the Stiefel manifold the easiest approach is to simulate $U_1(q \times r)$ with independent $N(0, 1)$ entries, and set $X_1 = U_1(U_1^T U_1)^{-1/2}$ using the symmetric square root of a positive definite matrix.

8.2 von Mises-Fisher distribution $\mathbf{F}_p(\kappa, \mu_0)$ on S_p

For general $p \geq 1$, the recommended method of simulation is an acceptance/rejection method due to Ulrich (1984), as modified by Wood (1994). This method uses a fractional linear transformation of a beta variate to provide an envelope for $u = x^T \mu_0$. It gives good efficiency across the whole range of values for κ . In particular, for large κ the distribution of $2(1 - u)$ is approximately the squared radial part of a multivariate normal distribution under the von Mises-Fisher model and of a multivariate Cauchy distribution under the envelope model, mimicking the efficiency calculations in (2.7).

Once the distribution of $u \in [0, 1]$ has been simulated, it is straightforward to whole von Mises-Fisher distribution by incorporating a uniformly distributed random direction y , say, on S_{p-1} (so y is a p -vector). More specifically, if $R = [R_1 \ \mu_0]$ is any $q \times q$ rotation matrix whose last column equals μ_0 , let $x = u\mu_0 + (1 - u^2)^{1/2} R_1^T y$.

For $p = 1$ the Ulrich-Wood method is essentially identical to the Best and Fisher (1979) method. One small exception to the recommendation to use the Ulrich-Wood method is the case $p = 2$ dimensions when u follows a truncated exponential

distribution and can be simulated more simply by the inverse method without any need for rejection (Fisher et al., 1987, p. 59).

8.3 Bingham distribution $\text{Bing}_p(A)$ on S_p or $\mathbb{R}P_p$

The BACG method developed in this paper is the first general-purpose acceptance/rejection simulation method for the Bingham distribution. However, earlier methods have been discussed in the literature for some special cases. In particular if $p = 1$, the BACG method reduces to Best and Fisher (1979) method for the von Mises distribution as noted in Section 6.

If $p = 2$ and either $0 = \lambda_1 < \lambda_2 = \lambda_3$ (bipolar case) or $0 = \lambda_1 = \lambda_2 < \lambda_3$ (girdle case), the simulation problem can be reduced to a one-dimensional problem. Best and Fisher (1986) developed effective envelopes in these cases, with efficiencies broadly comparable to the BACG method.

If the eigenvalues appear in pairs then the methods for the complex Bingham can be used. Kent et al. (2004) developed several simulation methods that sometimes are better than BACG.

The BACG method here supersedes the MCMC method of Kume and Walker (2006).

8.4 Fisher-Bingham distribution $\text{FB}(\kappa, \mu_0, A)$ on S_p

The Kent (or FB_5) distribution on S_2 is a special case of an aligned Fisher-Bingham distribution. An efficient simulation algorithm for FB_5 was developed by Kent and Hamelryck (2005). Otherwise, for other aligned FB distributions when $p = 2$ or for other values of p , the BACG-based method developed in Section 6 is the recommended method.

In particular, these methods supersede earlier acceptance-rejection methods developed by Wood (1987) for various special types of aligned Fisher-Bingham distribution on S_2 . They also supersede the acceptance rejection method of Scealy and Welsh (2011, Appendix A4) for a higher-dimensional version of the Kent distribution, for which the efficiency drops to 0 under high concentration when $p > 2$. In addition they supersede the MCMC method of Kume and Walker (2009) in the aligned case.

For non-aligned Fisher-Bingham distributions, it is difficult to make any firm theoretical statements about the behaviour of the algorithm in Section 6. However, under moderate concentration it is still likely to be preferable to the MCMC methods of Kume and Walker (2009).

8.5 Matrix Fisher distribution $\text{MF}(F)$ on $SO(r)$

When $r = 2$, $SO(2)$ is the same as S_1 and the matrix Fisher on $SO(2)$ is identical to the von Mises distribution on S_1 , so no new methodology is needed.

When $r = 3$ the accidental isomorphism in Section 5.3 reduces this case to the Bingham distribution on S_2 , which can be simulated by the BACG method.

Earlier methods to simulate the matrix Fisher distribution on $SO(3)$, now superseded by BACG, were based on MCMC algorithms. These include Green and Mardia (2006) and Habeck (2009).

The cases $r > 3$ are at least partly covered by the next subsection.

8.6 Matrix Fisher-Bingham distribution $\text{MFB}(F, A, C)$ on $\mathcal{V}_{r,q}$

For the general matrix Fisher-Bingham distribution $\text{MFB}(F, A, C)$ on $\mathcal{V}_{r,q}$, there is not yet a convenient and efficient A/R algorithm other than for the balanced matrix Bingham case, where a solution was given in Section 7. However, the recent MCMC algorithms of Hoff (2009) and Byrne and Girolami (2013) can deal with this case.

8.7 Product manifolds

Finally, the main setting not covered in this paper is the class of product manifolds where multivariate versions of directional models can be defined. There are a few special cases where acceptance rejection methods are available (e.g. Mardia et al. (2006, supplementary material) for the sine and cosine versions of a bivariate version of the von Mises distribution on the torus), but in general MCMC methods are needed.

References

- Best, D. J. and Fisher, N. I. (1979). Efficient simulation of the von Mises distribution, *J. Appl. Statist.* **28**: 152–157.
- Best, D. J. and Fisher, N. I. (1986). Goodness-of-fit and discordancy tests for samples from the Watson distribution on the sphere, *Austr. J. Statist.* **1**: 13–31.
- Byrne, S. and Girolami, M. (2013). Geodesic Monte Carlo on embedded manifolds, *Scandinavian Journal of Statistics* ??: to appear.
- Chikuse, Y. (2003). *Statistics on Special Manifolds*, Springer-Verlag, New York.
- Fisher, N. I., Lewis, T. and Embleton, B. J. J. (1987). *Statistical Analysis of Spherical Data*, Cambridge University Press, Cambridge.
- Ganeiber, A. M. (2012). *Estimation and simulation in directional and statistical shape models*, PhD thesis, University of Leeds.
- Green, P. J. and Mardia, K. V. (2006). Bayesian alignment using hierarchical models, with applications in protein bioinformatics, *Biometrika* **93**: 235–254.

- Habeck, M. (2009). Generation of three-dimensional random rotations in fitting and matching problems, *Comput. Statist.* **24**: 719–731.
- Hoff, P. D. (2009). Simulation of the matrix Bingham-von Mises-Fisher distribution with application to multivariate and relational data, *J. of Comput. and Graph. Stat.* **18**: 438–456.
- Kendall, D. G. (1984). Shape-manifolds, procrusten matrices and complex projective spaces, *Bull. London Math. Soc.* **16**: 81–121.
- Kent, J. T. (1994). The complex Bingham distribution and shape analysis, *J. Roy. Statist. Soc. B* **56**: 285–299.
- Kent, J. T., Constable, P. D. L. and Er, F. (2004). Simulation for the complex Bingham distribution, *Statistics and Computing* **14**: 53–57.
- Kent, J. T., Ganeiber, A. M. and Mardia (2012). Simulation of the bingham distribution in directional data analysis, in K. V. Mardia, A. Gusnanto, A. D. Riley and J. Voss (eds), *New Statistics and Modern Natural Sciences*, pp. 57–59.
- Kent, J. T. and Hamelryck, T. (2005). Using the Fisher-Bingham distribution in stochastic models for protein structure, in S. Barber, P. D. Baxter, K. V. Mardia and R. E. Walls (eds), *Proceedings in Quantitative Biology, Shape Analysis and Wavelets*, Leeds University Press, pp. 57–60.
- Kent, J. T., Mardia, K. V. and McDonnell, P. (2006). The complex Bingham quartic distribution and shape analysis, *J. Roy. Statist. Soc. B* **68**: 747–765.
- Kume, A. and Walker, S. G. (2006). Sampling from compositional and directional distributions, *Stat. Comput.* **16**: 261–265.
- Kume, A. and Walker, S. G. (2009). On the Fisher-Bingham distribution, *Stat. Comput.* **19**: 167–172.
- Mardia, K. V. and Jupp, P. E. (2000). *Directional Statistics*, Wiley, Chichester.
- Mardia, K. V., Taylor, C. C. and Subramaniam, G. K. (2006). Protein bioinformatics and mixtures of bivariate von Mises distributions for angular data, *Biometrics* **63**: 505–512.
- Scealy, J. L. and Welsh, A. H. (2011). Regression for compositional data by using distributions defined on the hypersphere, *J. Roy. Statist. Soc. B* **73**: 351–375.
- Ulrich, G. (1984). Computer generation of distributions on the m -sphere, *Appl. Statist.* **33**: 158–163.
- Wood, A. T. A. (1987). The simulation of distributions in the Fisher-Bingham family, *Communications in Statistics* **B16**: 885–898.

Wood, A. T. A. (1994). Simulation of the von Mises-Fisher distribution, *Communications in Statistics* **B23**: 157–164.